

Visualizations With Python (TLC Project)

December 5, 2025

```
[1]: #The purpose of this project is to conduct exploratory data analysis on a
      ↪provided data set.
#The goal is to clean data set and create a visualization.
```

```
#Part 1: Imports, links, and loading
#Part 2: Data Exploration and Data cleaning
#Part 3: Building visualizations
#Part 4: Evaluate and share results
```

```
[2]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import datetime as dt
import seaborn as sns
```

```
[3]: df=pd.read_csv('data/2017_Yellow_Taxi_Trip_Data.csv')
```

```
[4]: df.head(10)
```

```
[4]:   Unnamed: 0  VendorID      tpep_pickup_datetime      tpep_dropoff_datetime \
0    24870114        2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
1    35634249        1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
2   106203690       1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
3    38942136        2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
4    30841670        2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
5    23345809        2  03/25/2017 8:34:11 PM  03/25/2017 8:42:11 PM
6    37660487        2  05/03/2017 7:04:09 PM  05/03/2017 8:03:47 PM
7    69059411        2  08/15/2017 5:41:06 PM  08/15/2017 6:03:05 PM
8    8433159         2  02/04/2017 4:17:07 PM  02/04/2017 4:29:14 PM
9    95294817        1  11/10/2017 3:20:29 PM  11/10/2017 3:40:55 PM

  passenger_count     trip_distance  RatecodeID store_and_fwd_flag \
0              6            3.34          1             N
1              1            1.80          1             N
2              1            1.00          1             N
3              1            3.70          1             N
```

4	1	4.37	1	N			
5	6	2.30	1	N			
6	1	12.83	1	N			
7	1	2.98	1	N			
8	1	1.20	1	N			
9	1	1.60	1	N			
0	PULocationID	DOLocationID	payment_type	fare_amount	extra	mta_tax	\
0	100	231	1	13.0	0.0	0.5	
1	186	43	1	16.0	0.0	0.5	
2	262	236	1	6.5	0.0	0.5	
3	188	97	1	20.5	0.0	0.5	
4	4	112	2	16.5	0.5	0.5	
5	161	236	1	9.0	0.5	0.5	
6	79	241	1	47.5	1.0	0.5	
7	237	114	1	16.0	1.0	0.5	
8	234	249	2	9.0	0.0	0.5	
9	239	237	1	13.0	0.0	0.5	
0	tip_amount	tolls_amount	improvement_surcharge	total_amount			
0	2.76	0.0	0.3	16.56			
1	4.00	0.0	0.3	20.80			
2	1.45	0.0	0.3	8.75			
3	6.39	0.0	0.3	27.69			
4	0.00	0.0	0.3	17.80			
5	2.06	0.0	0.3	12.36			
6	9.86	0.0	0.3	59.16			
7	1.78	0.0	0.3	19.58			
8	0.00	0.0	0.3	9.80			
9	2.75	0.0	0.3	16.55			

[5]: df.size

[5]: 408582

[6]: df.describe()

count	Unnamed: 0	VendorID	passenger_count	trip_distance	\
mean	2.269900e+04	22699.000000	22699.000000	22699.000000	
std	5.675849e+07	1.556236	1.642319	2.913313	
min	3.274493e+07	0.496838	1.285231	3.653171	
25%	1.212700e+04	1.000000	0.000000	0.000000	
50%	2.852056e+07	1.000000	1.000000	0.990000	
75%	5.673150e+07	2.000000	1.000000	1.610000	
max	8.537452e+07	2.000000	2.000000	3.060000	

```

      RatecodeID PULocationID DOLocationID payment_type fare_amount \
count 22699.000000 22699.000000 22699.000000 22699.000000 22699.000000
mean 1.043394 162.412353 161.527997 1.336887 13.026629
std 0.708391 66.633373 70.139691 0.496211 13.243791
min 1.000000 1.000000 1.000000 1.000000 -120.000000
25% 1.000000 114.000000 112.000000 1.000000 6.500000
50% 1.000000 162.000000 162.000000 1.000000 9.500000
75% 1.000000 233.000000 233.000000 2.000000 14.500000
max 99.000000 265.000000 265.000000 4.000000 999.990000

      extra mta_tax tip_amount tolls_amount \
count 22699.000000 22699.000000 22699.000000 22699.000000
mean 0.333275 0.497445 1.835781 0.312542
std 0.463097 0.039465 2.800626 1.399212
min -1.000000 -0.500000 0.000000 0.000000
25% 0.000000 0.500000 0.000000 0.000000
50% 0.000000 0.500000 1.350000 0.000000
75% 0.500000 0.500000 2.450000 0.000000
max 4.500000 0.500000 200.000000 19.100000

      improvement_surcharge total_amount
count 22699.000000 22699.000000
mean 0.299551 16.310502
std 0.015673 16.097295
min -0.300000 -120.300000
25% 0.300000 8.750000
50% 0.300000 11.800000
75% 0.300000 17.800000
max 0.300000 1200.290000

```

[7]: df.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        22699 non-null   int64  
 1   VendorID         22699 non-null   int64  
 2   tpep_pickup_datetime  22699 non-null   object  
 3   tpep_dropoff_datetime 22699 non-null   object  
 4   passenger_count    22699 non-null   int64  
 5   trip_distance       22699 non-null   float64 
 6   RatecodeID          22699 non-null   int64  
 7   store_and_fwd_flag   22699 non-null   object  
 8   PULocationID        22699 non-null   int64  
 9   DOLocationID        22699 non-null   int64  

```

```

10 payment_type           22699 non-null   int64
11 fare_amount            22699 non-null   float64
12 extra                  22699 non-null   float64
13 mta_tax                22699 non-null   float64
14 tip_amount              22699 non-null   float64
15 tolls_amount            22699 non-null   float64
16 improvement_surcharge   22699 non-null   float64
17 total_amount             22699 non-null   float64
dtypes: float64(8), int64(7), object(3)
memory usage: 3.1+ MB

```

[8]: *#A box plot will be helpful to determine outliers and where the bulk of the data points reside in terms of trip_distance, duration, and total_amount*

#A scatter plot will be helpful to visualize the trends and patterns and outliers of critical variables, such as trip_distance and total_amount

#A bar chart will help determine average number of trips per month, weekday, weekend, etc.

[9]: *# Convert data columns to datetime*

```

df['tpep_pickup_datetime']=pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime']=pd.to_datetime(df['tpep_dropoff_datetime'])


```

[10]: *# Create box plot of trip_distance*

```

plt.figure(figsize=(7,2))
plt.title('trip_distance')
sns.boxplot(data=None, x=df['trip_distance'], fliersize=1);

```




[11]: *# Create histogram of trip_distance*

```

plt.figure(figsize=(10,5))


```

```
sns.histplot(df['trip_distance'], bins=range(0,26,1))
plt.title('Trip distance histogram');
```




[12]: #The majority of trips were journeys of less than two miles. The number of trips falls away steeply as the distance traveled increases beyond two miles.

```
[13]: # Create box plot of total_amount
plt.figure(figsize=(7,2))
plt.title('total_amount')
sns.boxplot(x=df['total_amount'], fliersize=1);
```




```
[14]: # Create histogram of tip_amount
plt.figure(figsize=(12,6))
ax = sns.histplot(df['tip_amount'], bins=range(0,21,1))
ax.set_xticks(range(0,21,2))
ax.set_xticklabels(range(0,21,2))
plt.title('Tip amount histogram');
```




[15]: #The distribution for tip amount is right-skewed, with nearly all the tips in the \$0-3 range.

```
[16]: # Create histogram of tip_amount by vendor
plt.figure(figsize=(12,7))
ax = sns.histplot(data=df, x='tip_amount', bins=range(0,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(0,21,1))
ax.set_xticklabels(range(0,21,1))
plt.title('Tip amount by vendor histogram');
```



[17]: #Separating the tip amount by vendor reveals that there are no noticeable → aberrations in the distribution of tips between the two vendors in the → dataset. Vendor two has a slightly higher share of the rides, and this → proportion is approximately maintained for all tip amounts.

```
[18]: # Create histogram of tip_amount by vendor for tips > $10
tips_over_ten = df[df['tip_amount'] > 10]
plt.figure(figsize=(12,7))
ax = sns.histplot(data=tips_over_ten, x='tip_amount', bins=range(10,21,1),
                  hue='VendorID',
                  multiple='stack',
                  palette='pastel')
ax.set_xticks(range(10,21,1))
ax.set_xticklabels(range(10,21,1))
plt.title('Tip amount by vendor histogram');
```



[19]: #The proportions are maintained even at these higher tip amounts, with the
 ↪exception being at highest extremity, but this is not noteworthy due to the
 ↪low sample size at these tip amounts.

[20]: df['passenger_count'].value_counts()

[20]:

1	16117
2	3305
5	1143
3	953
6	693
4	455
0	33


Name: passenger_count, dtype: int64

[21]: # Nearly two thirds of the rides were single occupancy, though there were still
 ↪nearly 700 rides with as many as six passengers. Also, there are 33 rides
 ↪with an occupancy count of zero, which doesn't make sense. These would
 ↪likely be dropped unless a reasonable explanation can be found for them.

[22]: # Calculate mean tips by passenger_count
`mean_tips_by_passenger_count = df.groupby(['passenger_count']).
 ↪mean()['tip_amount'])
 mean_tips_by_passenger_count`

```
[22]:          tip_amount  
passenger_count  
0              2.135758  
1              1.848920  
2              1.856378  
3              1.716768  
4              1.530264  
5              1.873185  
6              1.720260
```

```
[23]: # Create bar plot for mean tips by passenger count  
data = mean_tips_by_passenger_count.tail(-1)  
pal = sns.color_palette("Greens_d", len(data))  
rank = data['tip_amount'].argsort().argsort()  
plt.figure(figsize=(12,7))  
ax = sns.barplot(x=data.index,  
                  y=data['tip_amount'],  
                  palette=np.array(pal[::-1])[rank])  
ax.axhline(df['tip_amount'].mean(), ls='--', color='red', label='global mean')  
ax.legend()  
plt.title('Mean tip amount by passenger count', fontsize=16);
```



```
[24] :
```

#Mean tip amount varies very little by passenger count. Although it does drop
→noticeably for four-passenger rides, it's expected that there would be a
→higher degree of fluctuation because rides with four passengers were the
→least plentiful in the dataset (aside from rides with zero passengers).

[25]: # Create a month column
df['month'] = df['tpep_pickup_datetime'].dt.month_name()
Create a day column
df['day'] = df['tpep_pickup_datetime'].dt.day_name()

[26]: # Get total number of rides for each month
monthly_rides = df['month'].value_counts()
monthly_rides

[26]:

March	2049
October	2027
April	2019
May	2013
January	1997
June	1964
December	1863
November	1843
February	1769
September	1734
August	1724
July	1697

Name: month, dtype: int64

[28]: #The months are out of order.
#Reorder the results to put the months in calendar order.

[29]: # Reorder the monthly ride list so months go in order
month_order = ['January', 'February', 'March', 'April', 'May', 'June', 'July',
'August', 'September', 'October', 'November', 'December']

monthly_rides = monthly_rides.reindex(index=month_order)
monthly_rides

[29]:


January	1997
February	1769
March	2049
April	2019
May	2013
June	1964
July	1697
August	1724
September	1734

```
October      2027
November    1843
December    1863
Name: month, dtype: int64
```

```
[30]: # Show the index
monthly_rides.index
```

```
[30]: Index(['January', 'February', 'March', 'April', 'May', 'June', 'July',
           'August', 'September', 'October', 'November', 'December'],
           dtype='object')
```

```
[31]: # Create a bar plot of total rides per month
plt.figure(figsize=(12,7))
ax = sns.barplot(x=monthly_rides.index, y=monthly_rides)
ax.set_xticklabels(month_order)
plt.title('Ride count by month', fontsize=16);
```




```
[32]: #Monthly rides are fairly consistent, with notable dips in the summer months of July, August, and September, and also in February.
```

```
[33]: # Repeat the above process, this time for rides by day
daily_rides = df['day'].value_counts()
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
             'Saturday', 'Sunday']
```

```
daily_rides = daily_rides.reindex(index=day_order)
daily_rides
```

```
[33]: Monday      2931
Tuesday     3198
Wednesday   3390
Thursday    3402
Friday      3413
Saturday    3367
Sunday      2998
Name: day, dtype: int64
```

```
[34]: # Create bar plot for ride count by day
plt.figure(figsize=(12,7))
ax = sns.barplot(x=daily_rides.index, y=daily_rides)
ax.set_xticklabels(day_order)
ax.set_ylabel('Count')
plt.title('Ride count by day', fontsize=16);
```



```
[35]: #Surprisingly, Wednesday through Saturday had the highest number of daily rides,
→while Sunday and Monday had the least.
```

```
[36]: # Repeat the process, this time for total revenue by day
day_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
→'Saturday', 'Sunday']
```

```

total_amount_day = df.groupby('day').sum()[['total_amount']]
total_amount_day = total_amount_day.reindex(index=day_order)
total_amount_day

```

[36]:


day	total_amount
Monday	49574.37
Tuesday	52527.14
Wednesday	55310.47
Thursday	57181.91
Friday	55818.74
Saturday	51195.40
Sunday	48624.06

[37]:

```

# Create bar plot of total revenue by day
plt.figure(figsize=(12,7))
ax = sns.barplot(x=total_amount_day.index, y=total_amount_day['total_amount'])
ax.set_xticklabels(day_order)
ax.set_ylabel('Revenue (USD)')
plt.title('Total revenue by day', fontsize=16);

```




[38]:

#Thursday had the highest gross revenue of all days, and Sunday and Monday had ↴ the least. Interestingly, although Saturday had only 35 fewer rides than ↴ Thursday, its gross revenue was ~\$6,000 less than Thursday's-more than a 10% ↴ drop.

```
[39]: # Repeat the process, this time for total revenue by month
total_amount_month = df.groupby('month').sum()[['total_amount']]
total_amount_month = total_amount_month.reindex(index=month_order)
total_amount_month
```

```
[39]:      total_amount
month
January      31735.25
February     28937.89
March        33085.89
April         32012.54
May          33828.58
June         32920.52
July         26617.64
August        27759.56
September    28206.38
October       33065.83
November     30800.44
December     31261.57
```

```
[40]: # Create a bar plot of total revenue by month
plt.figure(figsize=(12,7))
ax = sns.barplot(x=total_amount_month.index, u
                  ↪y=total_amount_month['total_amount'])
plt.title('Total revenue by month', fontsize=16);
```



```
[41]: #Monthly revenue generally follows the pattern of monthly rides, with
    ↪noticeable dips in the summer months of July, August, and September, and
    ↪also one in February.
```

```
[42]: # Get number of unique drop-off location IDs
df['DOLocationID'].nunique()
```

```
[42]: 216
```


```
[43]: # Calculate the mean trip distance for each drop-off location
distance_by_dropoff = df.groupby('DOLocationID').mean()[['trip_distance']]

# Sort the results in descending order by mean trip distance
distance_by_dropoff = distance_by_dropoff.sort_values(by='trip_distance')
distance_by_dropoff
```

```
[43]:          trip_distance
DOLocationID
207           1.200000
193           1.390556
237           1.555494
234           1.727806
137           1.818852
...
      ...
51            17.310000
11            17.945000
210           20.500000
29            21.650000
23            24.275000
```

```
[216 rows x 1 columns]
```

```
[44]: # Create a bar plot of mean trip distances by drop-off location in ascending
    ↪order by distance
plt.figure(figsize=(14,6))
ax = sns.barplot(x=distance_by_dropoff.index,
                  y=distance_by_dropoff['trip_distance'],
                  order=distance_by_dropoff.index)
ax.set_xticklabels([])
ax.set_xticks([])
plt.title('Mean trip distance by drop-off location', fontsize=16);
```



[45]: #This plot presents a characteristic curve related to the cumulative density function of a normal distribution. In other words, it indicates that the drop-off points are relatively evenly distributed over the terrain. This is good to know, because geographic coordinates were not included in this dataset, so there was no obvious way to test for the distribution of locations.

#To confirm this conclusion, consider the following experiment:

```
#Create a sample of coordinates from a normal distribution-in this case 1,500 pairs of points from a normal distribution with a mean of 10 and a standard deviation of 5
#Calculate the distance between each pair of coordinates
#Group the coordinates by endpoint and calculate the mean distance between that endpoint and all other points it was paired with
#Plot the mean distance for each unique endpoint
```

[46]: # 1. Generate random points on a 2D plane from a normal distribution

```
test = np.round(np.random.normal(10, 5, (3000, 2)), 1)
midway = int(len(test)/2) # Calculate midpoint of the array of coordinates
start = test[:midway] # Isolate first half of array ("pick-up locations")
end = test[midway:] # Isolate second half of array ("drop-off locations")
```

2. Calculate Euclidean distances between points in first half and second half of array

```
distances = (start - end)**2
distances = distances.sum(axis=-1)
distances = np.sqrt(distances)
```


3. Group the coordinates by "drop-off location", compute mean distance

```

test_df = pd.DataFrame({'start': [tuple(x) for x in start.tolist()],
                       'end': [tuple(x) for x in end.tolist()],
                       'distance': distances})
data = test_df[['end', 'distance']].groupby('end').mean()
data = data.sort_values(by='distance')

# 4. Plot the mean distance between each endpoint ("drop-off location") and all other points it connected to
plt.figure(figsize=(14,6))
ax = sns.barplot(x=data.index,
                  y=data['distance'],
                  order=data.index)
ax.set_xticklabels([])
ax.set_xticks([])
ax.set_xlabel('Endpoint')
ax.set_ylabel('Mean distance to all other points')
ax.set_title('Mean distance between points taken randomly from normal distribution');

```




[47]: # The curve described by this graph is nearly identical to that of the mean distance traveled by each taxi ride to each drop-off location. This reveals that the drop-off locations in the taxi dataset are evenly distributed geographically. Note, however, that this does not mean that there was an even distribution of rides to each drop-off point.

[48]: # Check if all drop-off locations are consecutively numbered
df['DOLocationID'].max() - len(set(df['DOLocationID']))

[48]: 49

```
[49]: #There are 49 numbers that do not represent a drop-off location.  
  
#To eliminate the spaces in the histogram that these missing numbers would  
→create, sort the unique drop-off location values, then convert them to  
→strings. This will make the histplot function display all bars directly next  
→to each other.
```

```
[50]: plt.figure(figsize=(16,4))  
# DOLocationID column is numeric, so sort in ascending order  
sorted_dropoffs = df['DOLocationID'].sort_values()  
# Convert to string  
sorted_dropoffs = sorted_dropoffs.astype('str')  
# Plot  
sns.histplot(sorted_dropoffs, bins=range(0, df['DOLocationID'].max()+1, 1))  
plt.xticks([])  
plt.xlabel('Drop-off locations')  
plt.title('Histogram of rides by drop-off location', fontsize=16);
```



```
[51]: #Notice that out of the 200+ drop-off locations, a disproportionate number of  
→locations receive the majority of the traffic, while all the rest get  
→relatively few trips. It's likely that these high-traffic locations are near  
→popular tourist attractions
```

```
[ ]:
```